

Agentic reliability gets the right eval

Benchmarks have always defined the tempo of technological advancement. For a long time, agents were evaluated by whether they could succeed once. The old $\text{pass}@k$ benchmark, which measures single-trial success, picks that one winner from one game. But it doesn't capture dependability or answer: Can I rely on this every time this situation comes up?

The industry is converging on a better framework: pass^k , which was first [introduced](#) as part of the τ -bench benchmark for LLM-based agents in 2024. Pass^k refers to the probability that an agent succeeds every time across k trials. As k increases, pass^k drops. For example, an agent with a 75% per-trial success rate ($\text{pass}@k$) has only a ~42% chance of succeeding three times in a row. That is, its pass^k is just $(0.75)^3 \approx 42\%$.

For many customer-facing agents, repeatability is critical. Humans don't tolerate "mostly works" in real-life tasks. If an agent can't reliably reproduce successful behavior – if it still needs babysitting – its value to the customer collapses fast.

Desktop agents are becoming usable

A huge amount of SaaS user interface today is not about work but friction: excessive clicking and brittle abstractions that don't match what the user is trying to do at the moment. That's why agentic companies rush to automate desktop workflows. But in 2025, your experience might be this: clicking a few times was still faster and clearer than typing or speaking a verbose command to an LLM, waiting for a response, then iterating. It was unsurprising to see headlines declaring how AI tools actually slowed workers down.

But things are changing rapidly: If last year's computer-using agents were toddlers, able to take a few

But things are changing rapidly. If last year's computer-using agents were toddlers – able to take a few steps but constantly at risk of breaking something – this year they feel more like five-year-olds. They're still limited and can't handle deeply creative or ambiguous work. But they can walk steadily. They can follow instructions. And critically, they can repeat tasks that don't require heavy reasoning – like a child repeating words after their parents – reaching new milestones on pass^k.

As pass^k keeps improving, we will see the viable, cross-OS desktop agent that completes end-to-end tasks without constant human intervention. Capability growth compounds once reliability crosses a threshold. Eventually, computer-using agents that can see and operate an interface will sit on top of SaaS, which is essentially an automated human-defined workflow with a modern UI.

Hardware will simplify once humans no longer operate

The third trend ties directly to our company's vision: the autonomous computer company.

Most modern hardware is designed around human ergonomics. Apple perfected the trackpad because humans needed it. But if AI becomes the primary operator, moving, clicking and typing disappear. As computer-operating agents become more powerful, the hardware that hosts them will become less complicated. Agents eliminate unnecessary human-computer interaction. Humans issue intent. AI does the work.

This is why our end game from the first day isn't restricted to software. Over time, computers will be designed for agents first – whatever form that ultimately takes. We're entering a phase where AI hardware won't just talk (think Alexa). It will also do.

There's understandable anxiety around AGI and social disruption. For a period, agents will do a lot of what people do today. These concerns deserve serious attention.

But history suggests that technological shifts, while disruptive, have always created new types of work. The assembly line didn't end manufacturing jobs – it created entire new industries. When agents become truly reliable, the challenge may shift to an undersupply of human labor for the new problems that still require human judgment and creativity.

What makes 2026 different is that we're crossing a threshold. Desktop agents are moving from research labs to production tools. They're becoming reliable enough that businesses will start depending on them. The question is no longer whether agents can do what humans do; it's how we design human-AI collaboration patterns that make this transition smooth.

2026 will be the year when desktop agents stop being a toy and start doing actual human work. And this is just the beginning.

Building autonomous computers doesn't mean replacing humans. It means cooperation.

Free your hands from the computer. Download Simular today for free.

Try Simular ↗

