

Authors:



Inna Lebedeva



Sergii Opanasenko



Kateryna Reshetilo



5

15.01.2026 ● 20 min read

Technologies

AI agents are no longer science fiction sidekicks - they're slipping quietly into spreadsheets, call centers, and CRMs. What began as experiments in automation is now reshaping how businesses operate day to day.

To map this shift, we analyzed 542 AI agent development jobs posted on Upwork - one of the world's largest freelance marketplaces. This dataset gives a rare, bottom-up view of what businesses are actually building and paying for in 2026. The picture that emerges is less about flashy demos and more about real infrastructure - agents that handle the grunt work, amplify human teams, and are fast becoming a standard feature of modern business.

And we're not just observers: as a Top Rated Upwork agency with a long track record in AI development, we see these trends firsthand in client projects, giving us a unique vantage point to speak about where agent technology is headed.

- **Python rules:** 52% of projects rely on it; LangChain and Pinecone anchor the stack.
- **Practical wins lead:** Back-office automation (15.2%) and customer support (14.8%) top use cases.
- **Marketing drives demand:** 17.6% of projects, pulling outreach, content, and support agents with it.
- **Costs cluster:** Most jobs fall in the \$20-60/hour range for 1-3 months, but extremes stretch from \$5 to \$600/hour.
- **From hype to habit:** AI agents are shifting from experiments to a standard layer of business operations.

请更新你的浏览器

你使用的似乎是旧版的或不支持的浏览器。要继续使用Canva，请更新至以下任意一种浏览器的最新版本：




Chrome



Firefox

Safari
(仅限 macOS)

Edge






The tech stack behind AI agents shows how the industry is shifting from prototypes to production. We analyzed 542 job posts, pulling mentions of languages, frameworks, databases, model providers, and tools. The pattern is clear: a few technologies dominate, challengers emerge in niches, and legacy tools fade. What follows is the breakdown of what's shaping AI agents in 2026 - and why it matters for startups, vendors, and the ecosystem.

225 out of 542 jobs specified their programming language preferences. If AI agents are the new apps, Python is their operating system. In over half of the projects (52%), Python sits at the core - more than all other languages combined. It has become the lingua franca of AI development, the environment where TensorFlow, PyTorch, LangChain, and Hugging Face live. Startups turn to it for fast prototyping; enterprises rely on it for its deep ecosystem and production-ready stability.

TOP PROGRAMMING LANGUAGES FOR AI AGENT



greenice

225 of 542 AI agent jobs specified languages - and in 52% of them, Python ruled. Original research by Greenice.

Once the “brain” of an agent is built in Python, other languages take over to make it usable at scale. Node.js (17%) and Go (12%) are the muscle behind real-time APIs and concurrent workloads. They turn Python’s intelligence into reliable services, keeping latency low and throughput high.


On the user-facing side, JavaScript (10%) and TypeScript (6%) act as connectors. They bring AI logic into web apps, dashboards, and client integrations, smoothing the bridge between intelligence and interaction.

At the margins, Java, Rust, and PHP (around 1% each) show up in pockets of experimentation or legacy contexts. Their presence is too small to shape the ecosystem - and the absence of Swift and R underlines a larger shift: mobile-native and traditional data science stacks are no longer driving innovation in this space.

For startups, the path is simplified: Python talent comes first, but expect to complement it with Node.js or Go for backends and TypeScript for interfaces. For vendors, the implication is clear: prioritize SDKs and integrations for this triad - Python, Node.js, and Go - if you want to reach across the agent economy.

Every AI agent needs memory - not for birthdays, but for facts and fuzzy recall. Out of 542 jobs we analyzed, 133 mentioned databases or vector stores, revealing a split between AI-native tools like Pinecone, Weaviate, and Qdrant, and traditional backbones like Postgres adapting to the new era. The result is less a battlefield than a patchwork ecosystem still deciding whether memory should be specialized or standardized.

TOP DATABASES AND VECTOR STORES FOR AI AGENTS



Out of 542 AI agent jobs, 133 mentioned databases or vector stores - from Pinecone to Postgres. Original research by Greenice.

At the top of the stack sits Pinecone (22.6%), which has quietly become the AWS of recall. It abstracts away the pain of managing clusters and plugs neatly into frameworks like LangChain - making it the default for startups racing to ship retrieval-augmented generation (RAG) systems. If Pinecone is the “cloud of memory,” then Weaviate (16.5%), Qdrant (4.5%), and Milvus (4.5%) are the open-source insurgents - labs where developers roll up their sleeves to trade SaaS convenience for control and cost savings. For indie hackers and quick prototypes, Chroma (0.8%) is the garage-built kit: cheap, scrappy, good enough to get ideas off the ground.

Then there’s the steel frame of enterprise software: PostgreSQL (18.8%). With its pgvector extension, Postgres has been retrofitted for the AI era without losing its reputation for reliability. MySQL (6.8%), meanwhile, clings to legacy projects, while Redis (8.3%) and MongoDB (4.5%) offer speed and schema flexibility - now with vector search bolted on. The message from the old guard is clear: why add another database if your existing one can do the job?


In the middle ground are tools like Faiss (9.8%), Facebook’s open-source performance engine, favored by teams who want raw speed over ease of use. Elasticsearch (1.5%) - once the search darling - is fading into niche use, while SQLite (1.5%) survives in lightweight prototypes and embedded agents.

For startups, the choice is existential: do you buy Pinecone’s convenience and risk lock-in, or bet on an open-source stack and swallow the DevOps pain? For vendors, the writing is on the wall: integrations with Pinecone and Weaviate aren’t optional anymore.

Every AI agent needs a brain, and the choice of that brain is the most strategic decision a builder can make. Go with the proven

standard and you move fast; bet on a challenger and you may gain control, cost savings, or enterprise trust. Out of 542 jobs we analyzed, 366 mentioned LLMs - proof that model selection isn't an afterthought, but the foundation on which the rest of the agent stack sits.

TOP LLMs & MODEL PROVIDERS FOR AI AGENTS



greenice

366 of 542 AI agent jobs mentioned LLMs - the brains powering the stack.
Original research by Greenice.

For now, the gravitational pull of OpenAI is undeniable. With more than 70% of projects naming its APIs, GPT has become the default way to give an agent reasoning power. Developers flock to it because it works, integrates easily, and carries the weight of a maturing ecosystem. But speed comes at a price - both in dollars and in dependence on a single provider.


Claude, at nearly 16,6% of mentions, has become the safe bet for enterprises, prized for its careful alignment and guardrails. Google's Gemini, though smaller in share, carries the promise of tight integration with the company's cloud and productivity stack. And then there are the open-source contenders - Mistral and Meta's Llama - smaller in adoption but important for what they represent: local control, cost efficiency, and the freedom to self-host. Hugging Face underpins much of this open-source ecosystem, acting as the distribution hub where models are discovered, shared, and fine-tuned.

What's striking is that this isn't a winner-take-all race. Many serious teams are blending models rather than betting on just one. OpenAI often anchors the system, but it's paired with Claude for sensitive enterprise contexts, or with Llama and Mistral for cheaper batch tasks. The smartest builders design for flexibility, keeping the door open to multiple providers as the landscape shifts.

For entrepreneurs, the message is simple but strategic: OpenAI will get you to market fastest, but building with challengers can give you cost advantages, data control, or enterprise trust. For vendors, the implication is clear: multi-model support is now table stakes. The brains of the agent economy may still be overwhelmingly GPT, but the systems around it are already preparing for a plural future.

Once, the hard part of AI was training the models. Today, the frontier has shifted: the real challenge is engineering agents - wiring together models, memory, and logic so they behave like useful systems. Out of 542 AI agent jobs we analyzed, 126 mentioned frameworks or tooling. And the data shows something striking: there's now a dedicated stack for agent engineering, distinct from the old world of machine learning.

TOP AI ENGINEERING FRAMEWORKS & TOOLS FOR AI AGENTS



126 of 542 AI agent jobs cited frameworks or tooling 5 - agent engineering now has its own stack. Original research by Greenice.

LangChain (55.6%) towers above the rest. It's the Swiss Army knife of agent development - a framework that plugs large language models into memory, databases, APIs, and tools. If raw LLMs are the "brains," LangChain is the nervous system. Close behind, CrewAI (9.5%) and Autogen (5.6%) explore the next frontier: multi-agent coordination. Think of them as project managers for fleets of AI workers, orchestrating collaboration instead of solo intelligence.

Agents aren't useful unless they can access the right knowledge at the right time. LlamaIndex (7.1%) specializes in structuring enterprise data so AI can query it. Haystack lingers as a niche option, but the message is clear: retrieval is no longer an afterthought - it's a core layer of the stack.


The AI boom has been powered by cloud APIs, but there's a countercurrent: Ollama (4.0%). It lets developers run LLaMA and other open models locally, on laptops or servers. It's the privacy and cost-control play - important for startups who can't afford Pinecone-sized bills or for enterprises wary of sending data outside their walls.

Once the stars of the show, PyTorch (6.3%) and TensorFlow (5.6%) now play supporting roles in agent jobs - used for fine-tuning or custom model work, not for building systems from scratch. Pandas and Scikit-Learn still appear, but mostly at the edges, where old data science habits meet new agent workflows.

Agent engineering now has its own playbook. For startups, the shift lowers the barrier to entry: success isn't about training a model, but about stitching one into a working system. For vendors, the message is blunt: integrations with LangChain and LlamaIndex are no longer nice-to-haves - they're table stakes.

The AI agent boom isn't just happening in code editors. It's unfolding in drag-and-drop canvases, where workflows and data are stitched together visually. Out of 542 jobs we analyzed, nearly half (247) mentioned no-code or low-code tools. That's a strong signal: businesses don't just want smarter agents - they want faster ways to build them.

NO-CODE/LOW-CODE PLATFORMS FOR AI AGENTS



greenice

247 of 542 AI agent jobs mentioned no-code or low-code tools - showing demand for faster, visual ways to build. Original research by Greenice.

The core layer is orchestration: chaining LLM calls, APIs, and logic into usable systems. n8n (38.1%) leads as the open-source favorite for complex automations, where engineers still design but don't need to hand-code every step. Zapier (27.9%) is the opposite: plug-and-play for non-technical teams who just want to connect apps to AI services. Make.com (15.0%) sits in between, offering a visual, design-driven interface that appeals to product teams who value clarity as much as power. Together, these three account for more than 80% of all no-code mentions - the operating system for lightweight AI automation.


Beneath the workflows are lightweight databases where agents store context. Airtable (10.5%) provides structured storage - a spreadsheet that acts like a backend. Notion (4.0%) plays the unstructured role: a free-form knowledge base readable by both humans and machines.

For startups, no-code lowers the cost of experimentation: you can prove ideas fast before investing in full custom

development. For vendors, it's a roadmap for integrations: if your product doesn't plug into n8n, Zapier, or Make, you risk being left out of the workflows where adoption happens. And for engineers, the message is nuanced: no-code doesn't erase your role - it changes it, shifting you from raw coder to system designer.

AI agents aren't just chat windows anymore - they're becoming voices you can dial, listen to, and talk with. Out of 542 jobs we analyzed, 181 mentioned voice, speech, or audio tools. That's a clear signal: natural conversation is moving from novelty to expectation.

TOP VOICE, SPEECH & AUDIO TECH FOR AI AGENTS



greenice

181 of 542 AI agent jobs mentioned voice, speech, or audio tech - conversation is shifting from novelty to expectation. Original research by Greenice.

At the foundation is Twilio (23.2%), the backbone of telephony and call routing. It provides the pipes that turn a text-based agent into something you can actually phone.

Above that sit Vapi (16.6%) and Retell (13.3%), the new platforms built specifically for real-time AI agents. They handle the hard parts of voice interaction - low-latency streaming, back-and-forth dialog, and orchestrating the LLM in the loop. If Twilio is the phone line, Vapi and Retell are the call centers where the agent actually speaks.

For speech recognition, Whisper (12.2%) has emerged as the open-source standard. It's prized for multilingual transcription and its developer-friendly ecosystem. Traditional ASR systems (3.9%) still appear, often tied to legacy enterprise deployments, but Whisper is increasingly the default for builders who want flexibility.


On the other side of the loop is speech synthesis. ElevenLabs (14.4%) leads the new wave with synthetic voices that carry tone and emotion, making conversations sound human instead of robotic. More generic TTS systems (16.6%) still power budget-friendly or

utilitarian use cases, but ElevenLabs is quickly setting the benchmark for quality.

Voice agents are no longer end-to-end projects - they're stitched together from modular parts. Twilio provides the pipes, Whisper the ears, ElevenLabs the voice, and Vapi or Retell the conversation engine. For entrepreneurs, this lowers the barrier to launching voice-first products: you assemble components instead of reinventing them. For vendors, the priority is integration - the winners will be the ones that fit smoothly into this voice stack.

AI agents aren't chasing sci-fi dreams. They're doing something far less glamorous but far more valuable: automating the everyday work that keeps businesses running. From the job descriptions, we mapped 512 projects to specific use cases. The vast majority fall into practical categories that cut costs, save time, and boost productivity. Here are the five leading use cases - and why they matter.

TOP USE CASES FOR AI AGENTS



512 of 542 AI agent jobs mapped to concrete use cases - showing agents focus less on sci-fi dreams and more on cutting costs and boosting productivity. Original research by Greenice.

The most common use case for AI agents isn't customer service or creativity - it's paperwork. In our dataset, 15.2% of projects focused on back-office automation: agents that quietly handle the repetitive, error-prone tasks that keep businesses running.

N8n AI Agent Expert

Posted 3 weeks ago



Worldwide



Specialized profiles can help you better highlight your expertise when submitting proposals to jobs like these. [Create a specialized profile.](#)

Summary

I need an expert to help build an AI agent that I talk to then trigger things to happen within a third party SaaS app. I have a distribution company where i am trying to automate the following flows:

- Creating an invoice (Sales order) to be sent to client.
- Quoting a client new parts that they request a quote for
- Sending purchase orders to my vendors of the parts I need to purchase from them to resell to my clients.

Can discuss further over a call



\$200.00

Fixed-price



Expert

I am willing to pay higher rates for the most experienced freelancers

Back-office automation job example

Think of them as digital clerks: generating invoices, pushing purchase orders, juggling records, and syncing CRMs with ERPs. At the high end, they file compliance reports, manage translation pipelines, or coordinate with other agents. They don't talk, they don't dazzle - but they erase the grunt work that used to bury human teams.

Typical builds combine workflow platforms (n8n, Zapier, Make.com), AI frameworks (LangChain, CrewAI), LLMs (OpenAI, Claude), vector databases for RAG, OCR for documents, and integrations with tools like Airtable, HubSpot, or QuickBooks.

Back-office agents may lack glamour, but they deliver clear ROI. By automating paperwork-heavy workflows, they cut overhead, reduce mistakes, and often serve as a company's first step into AI adoption.

Support remains one of the fastest wins for AI agents. In our dataset, 14.8% of projects focused on customer-facing automation: systems that answer questions instantly, triage issues, and deliver 24/7 responsiveness. These agents don't just cut call-center costs - they reset customer expectations for speed and availability.

AI Engineer Needed for LLM-Powered Chatbot Development

Posted 3 weeks ago  Worldwide

 Specialized profiles can help you better highlight your expertise when submitting proposals to jobs like these. [Create a specialized profile.](#)

Summary

We are looking for an experienced AI Engineer to build a chatbot powered by Large Language Models (LLMs). The chatbot should assist with customer support, query handling, and intelligent automation tailored to our business use case. The system should be scalable, reliable, and able to integrate with existing APIs.

Key Requirements:

- Strong background in Artificial Intelligence & Machine Learning.
- Proven experience working with LLMs (GPT, LLaMA, Claude, Mistral, etc.).
- Expertise in chatbot development (custom NLP pipelines or LLM-based).
- Hands-on experience with frameworks such as LangChain, Hugging Face Transformers, or Rasa.
- Ability to fine-tune models and design prompt engineering strategies.
- Experience deploying solutions with Flask/FastAPI + Docker.
- Familiarity with vector databases (Pinecone, Weaviate, FAISS) for retrieval-based systems.

Deliverables:

- A working AI-powered chatbot integrated with our APIs.
- Ability to handle multi-turn conversations and maintain context.
- Admin interface or API endpoints to update and manage knowledge base.
- Deployment on cloud (AWS/GCP/Azure) with monitoring & scaling support.

Nice to Have:

- Experience with speech-to-text / text-to-speech integration.
- Knowledge of MLOps and production monitoring.
- Prior Upwork experience delivering AI/LLM projects.

Customer support & live chat job example

These agents triage emails and chats, pull answers from knowledge bases, and route tickets automatically. On voice lines they check orders, handle cancellations, book appointments, and cover FAQs. They work across every channel - WhatsApp, email, web, phone - and hand off smoothly when a human needs to step in.

Typical systems pair LLMs (OpenAI, Claude, LLaMA, Mistral) with agent frameworks like LangChain and workflow tools (n8n, Zapier, Make). They plug into CRMs (HubSpot, Salesforce, GoHighLevel), telephony stacks (Asterisk, Twilio, JustCall), and vector databases (Pinecone, Weaviate, FAISS) to power real-time answers.

Customer support is where people feel AI most directly. Done well, these agents slash response times, reduce costs, and boost satisfaction. They don't replace humans outright - they scale the routine so human teams can focus on the cases that matter most.

The phone call isn't dead - it's changing hands. With 7.6% of projects in our dataset, voice agents are taking over appointment booking, lead qualification, and customer verification once handled by call-center staff.



Voice agents listen with speech-to-text, think with LLMs, and reply with lifelike text-to-speech. They route calls, log transcripts, catch voicemails, and push data into CRMs or booking systems. The advanced ones juggle multiple languages and plug directly into

industry-specific software.

These agents combine telephony platforms (Twilio, Asterisk, JustCall) with voice AI services (Vapi, Retell, Whisper, ElevenLabs), orchestration tools (n8n, Make.com), and LLMs (OpenAI, Claude). They integrate with CRMs and booking software to close the loop.

Industries like real estate, healthcare, and finance still depend on phone calls. AI voice agents cut missed calls, extend availability to 24/7, and scale outbound sales or follow-ups. They don't replace humans, but they reliably cover the first line of conversation so no lead or inquiry slips through.

In sales, time is everything. With 7.4% of projects in our dataset, lead and outreach agents are emerging as force multipliers - scraping prospects, personalizing first contact, and letting human sales teams focus on closing instead of chasing.

Social Media Agent

Posted last month  Worldwide

 Specialized profiles can help you better highlight your expertise when submitting proposals to jobs like these. [Create a specialized profile.](#)

Summary

Build an automated system to identify, qualify, and reach out to social creators/influencers administrators for book promotion. The system will handle the complete outreach process from discovery to personalized email deployment.

Core Automation Workflow:

- Identify relevant influencers using Apify scraping
- Qualify and score prospects based on engagement and relevance
- Enrich prospect data with interests and recent posts
- Generate personalized email messages using AI
- Deploy email campaign

 **\$875.00**
Fixed-price

 **Intermediate**
I am looking for a mix of experience and value

Lead & outreach automation job example

They pull leads from social feeds, public data, or CRMs, then reach out over WhatsApp, SMS, email, or voice. They qualify replies, log everything into sales systems, and even personalize pitches - from LinkedIn intros to influencer outreach - so humans can step in only when it counts.

Typical setups blend workflow tools (n8n, Make.com, Zapier) with CRMs (HubSpot, Airtable, GoHighLevel), messaging APIs (WhatsApp Business, Twilio), and voice platforms (Retell, Vapi). LLMs like OpenAI and Claude generate personalized messages at scale.

Manual outreach is slow, inconsistent, and costly. AI outreach agents ensure every lead is contacted quickly and consistently,

qualifying prospects before humans step in. They don't replace sales reps - they sharpen their focus so the team spends time only where it counts.

Creativity is no longer the sole domain of humans. In our dataset, 7.2% of projects involved agents producing blog drafts, ad copy, images, or marketing assets. These systems sit at the intersection of automation and creativity, helping teams keep up with the relentless demand for fresh content.

AI Content Automation Expert (Existing N8N Workflow + Customisation)

Posted 3 weeks ago  Worldwide



Specialized profiles can help you better highlight your expertise when submitting proposals to jobs like these. [Create a specialized profile.](#)

Summary

We are looking for an experienced automation builder who already has a working AI-powered content creation workflow (preferably in N8N or similar) that can take us from trend/idea discovery → script generation → final post creation for:

LinkedIn (Video posts, carousels)

YouTube (video scripts/shorts)

X/Twitter (threads, short posts)

Website/Medium (SEO-informed articles or brand content)

We'd like to focus mainly on LinkedIn and YouTube video content to begin with, building and scaling our CEO's personal brand.

The ideal candidate will:

- Present the content automation system they have already built and proven to work.
- Adapt/customise it to our CEO's personal brand voice, content pillars, and target audience.
- Enable high-volume, high-quality content production with minimal manual steps.

Content & creative asset generation job example

They mine Reddit threads and news feeds for themes, draft LinkedIn posts or TikTok scripts, and generate visuals with Midjourney, Runway, or Pika. Then they push content through publishing APIs, manage calendars, and track performance. The advanced ones stitch entire video pipelines from ingest to edit to final render - no human editor required.

These agents rely on LLMs (OpenAI, Claude) with frameworks like LangChain or LangGraph, paired with media tools (Runway, Pika, D-ID, ElevenLabs) for visuals and voice. They use scraping tools (Apify, Playwright) for research, social/CMS integrations for distribution, and orchestration platforms (n8n, Zapier, Make) to hold it all together. Infrastructure often runs on cloud GPUs with compliance gates for regulated niches.

Modern publishing rewards speed, volume, and consistency. Content agents compress research, creation, and distribution into repeatable workflows, letting lean teams ship more - and more personalized - assets while freeing human creatives to focus where it matters most.

Beyond the top five use cases, AI agents fragment into niches. Some will stay small, others could become tomorrow’s defaults.

Qualify leads, draft outreach, and keep pipelines warm - freeing reps to focus on closing.

Parse datasets, generate reports, and flag anomalies in finance, ops, and research.


Book meetings, send reminders, and manage appointments across services and clinics.

Tutoring bots, recruiting assistants, finance helpers, even early AI doctors - experiments in fields where trust is still unproven.

The long tail is where new roles emerge. Sales, analyst, and scheduling agents are already gaining ground; the rest show where demand may tip next.

AI agents aren’t confined to labs or startups - they’re finding their way into every corner of the economy. By analyzing job descriptions, we identified industries for 522 out of 542 projects, revealing where demand is strongest and how adoption is spreading across sectors. The data shows both clear leaders and a diverse long tail of experimentation.

TOP INDUSTRIES FOR AI AGENTS



522 of 542 AI agent jobs tied to specific industries - revealing clear leaders and a long tail of experimentation. Original research by Greenice.

With 17.6% of projects, no industry uses AI agents more than marketing and sales. The reason is simple: customer acquisition is where speed and personalization matter most - and where automation delivers measurable results.

These agents spin up ad copy, blog posts, and social campaigns at scale, then chase leads across email, chat, or voice. They qualify prospects, auto-schedule meetings, and track performance with built-in analytics. In short, they act as a marketing department in a box - filling the funnel faster and cheaper than human teams ever could.

Typical builds pair workflow tools like n8n, Zapier, and Make with CRMs such as GoHighLevel and Airtable. Voice platforms (Twilio, Vapi, Retell) handle calls, while LLMs generate both persuasive copy and personalized conversations. Together, these form a “marketing stack in miniature” that blends outreach, content generation, and support into one automated pipeline.

Outreach and content agents are surging, driven by Marketing & Sales’ need to fill funnels fast. Support agents feed the same loop, qualifying leads and handling handoffs. For startups, it’s scale without extra headcount; for vendors, the demand is clear - CRMs, voice stacks, and content pipelines.

With 13.2% of projects, enterprise and business software is the second-largest industry adopting AI agents. Here, agents don’t usually stand alone - they live inside CRMs, spreadsheets, and productivity suites, transforming familiar tools into smarter copilots.

These copilots generate forms, update records, and sync data across systems. They double as support bots or knowledge assistants, pulling answers from company data. At the high end, they evolve into multi-agent SaaS frameworks with white-label options, turning AI into an embedded feature rather than a standalone product.

Typical deployments use LLMs (OpenAI, Claude, LLaMA, Mistral) with frameworks like LangChain, paired with vector databases (Pinecone, Weaviate, FAISS) for retrieval. They run on enterprise-ready infrastructure such as FastAPI, Docker, and cloud platforms, and often surface through React or Next.js UIs. Many also integrate with workflow orchestrators (n8n, Make, Zapier) to extend back-office automations deeper into corporate systems.

Enterprise AI isn’t about flashy chatbots - it’s about embedding intelligence into the tools companies already use. Back-office agents cut paperwork, knowledge copilots speed up search, and orchestrators link it all. For startups, it’s a chance to build agent-first features; for vendors, integration is now table stakes.

Clinics are turning to AI agents, from call-handling receptionists to symptom triage assistants. Compliance remains a hurdle, but momentum is clear: healthcare already makes up 8.1% of projects, showing how fast agents are becoming force multipliers for staff.

Voice agents book appointments, send reminders, and handle intake questions. Others cut no-shows by managing schedules, or streamline back-office work like inventory and vendor tracking. Knowledge copilots use RAG to pull reliable answers from medical records and trial data. The result: less paperwork for staff, more focus on patients.

These agents rely on voice AI platforms (Vapi, Retell, Twilio) for patient communication, workflow tools (n8n, Make, Zapier) for automation, and LLMs with RAG (OpenAI, LangChain, vector databases) for knowledge retrieval. Crucially, they integrate with EHR and compliance-ready APIs to meet HIPAA and GDPR standards.

Healthcare runs on efficiency and trust. AI agents handle admin, scheduling, and reminders, freeing clinicians for care - but they must stay compliant. For startups, the wins are in voice, scheduling, and back-office workflows; for vendors, it's about building regulation-ready integrations.

AI agents here aren't creative copilots but precise automators, handling transactions, reporting, and risk monitoring in workflows where mistakes are costly. With 6.5% of projects, finance is carving out a clear role: dependable, compliance-ready support.

These agents process invoices, reconcile accounts, and scan financial statements for anomalies. They generate KPI dashboards, draft performance reports, and, through voice integrations, handle client queries or schedule meetings for advisors. In short, they keep the books balanced and the conversations flowing.

Finance agents typically combine LLMs (OpenAI, Claude) with workflow orchestration (n8n, Zapier, Make) and vector databases (Pinecone, FAISS) for retrieval. They rely on OCR to process invoices, integrate with financial software (QuickBooks, Xero, Stripe), and often include voice tools (Twilio, JustCall) for client engagement - all with compliance as a non-negotiable requirement.

In finance, agents don't sell - they safeguard. They cut grunt work while keeping outputs consistent, accurate, and audit-ready. For startups, the play is building agents that pass the trust test; for vendors, it's integrating with financial systems and compliance frameworks.

Content never sleeps, and neither can the teams behind it. In media and publishing, AI agents draft copy, generate visuals, and curate feeds. At 5.9% of projects, they're not replacing creators but reshaping how stories and campaigns get produced at scale.

These agents scrape forums and news feeds, then spin summaries into news posts, marketing copy, or even long-form drafts. They manage eBook pipelines, edit transcripts, generate captions, and assemble videos with minimal human effort - compressing the publishing cycle from days to minutes.

Typical builds rely on LLMs (OpenAI, Claude, Perplexity, Groq) orchestrated with workflow tools (Zapier, Make, n8n). They plug into media AI platforms (Runway, Midjourney, Dify, Flowise), use OCR or image analysis for editing, and connect with publishing systems through Webflow, CMS APIs, or YouTube integrations.

Publishing rewards speed and volume. AI agents let media teams scale from social clips to full campaigns without adding headcount - not replacing creatives, but acting as essential assistants in the production room.

Beyond the top five industries, adoption spreads into a wider set of verticals. Some are steady adopters embedding agents into daily workflows, others are cautiously experimenting, and a few are wildcards chasing niche but high-potential ideas.

Real estate (4.8%) leans on agents for listings, scheduling showings, and lead follow-up tasks. Education (3.1%) experiments with tutoring bots and study copilots. eCommerce (2.9%) uses agents for back-office tasks like inventory and order processing.

Travel and hospitality (2.5%) test agents for booking and customer support, areas that overlap with Voice & Call Automation and Support Agents. Home services (2.3%) apply similar models, using agents to manage appointments, quotes, and service requests.

Entertainment and gaming (1.9%) push into experimental territory with AI characters and creative copilots. Manufacturing (1.5%) looks at workflow monitoring and compliance automation, mirroring back-office use cases. HR (1.5%) adopts recruiting and scheduling assistants, while legal (1.3%) tests agents for policy Q&A and document review. These may be small today, but they highlight how AI is probing every corner of the economy, often starting in niches before scaling outward.


While the biggest gains are in marketing, support, and back-office tasks, the long tail shows AI agents are seeping into every corner of the economy - starting small in niche roles, but hinting at much broader adoption ahead.

"How much does it cost to build an AI agent?" Founders ask it first, but there's no simple answer. Costs swing from a basic invoicing bot to a healthcare voice agent tied into regulated systems. Still, freelance job postings reveal what buyers expect - in hourly rates and timelines.

The bulk of postings cluster in the \$15-60/hour range (32% of jobs), with \$20-40/hour the most common bracket. This reflects the global freelance market, where experienced solo developers or small teams handle mid-complexity projects - from content agents drafting copy to outreach bots wiring into CRMs.

At the extremes, the story changes. Lowest rates (from \$5/hour) are often tied to projects that sound simple - scraping influencers or spinning up basic chatbots - but in reality are underestimated by clients. These figures are less market price than placeholder, signaling exploratory budgets.

TOP HOURLY RATES FOR AI AGENT DEVELOPMENT




229 jobs named rates: most fall at \$20-40/hour, outliers run from \$5 placeholders to \$600 high-stakes builds. Avg: \$35.08. Original research by Greenice.

On the other end, highest rates (up to \$600/hour) appear in mission-critical builds: voice receptionists in healthcare, enterprise automation copilots, or multi-agent orchestration in financial systems. Here, companies aren't just buying hours - they're buying confidence, risk management, and integration expertise. The weighted average - \$35.08/hour - lands neatly within Upwork's broader benchmarks for AI engineering talent (\$30-65/hour).

Most clients imagine AI projects as short-term experiments: 1-3 months with less than 30 hours per week (36% of jobs). This is the sweet spot for MVPs - like a scheduling agent for a clinic or a back-office bot to process invoices. Another 17.5% expected turnaround in under a month, reflecting rapid prototypes or low-code configurations where speed outweighs depth.

The heavier commitments - 30+ hours per week over 1-3 months (13.5%) or 6+ months (13.5%) - point to larger-scale builds: enterprise copilots, healthcare workflow agents, or integrations spanning CRMs, ERPs, and compliance systems. A small share (4%) landed in the mid-scale range of 3-6 months, while 28% didn't specify a timeline at all, underscoring that many buyers simply don't know how long AI agents should take - leaving scoping to the developers.

TIME TO DEVELOP AN AI AGENT



383 out of 542 jobs set timelines: most expect 1-3 month MVPs, few commit long-term, and 28% give no estimate. Original research by Greenice.

Read Also:

[AI agent development cost 2026: Expectation vs reality](#)

The market is still feeling its way. At the low end, underpriced projects reflect misunderstanding of technical and compliance hurdles. At the high end, buyers pay a premium for reliability in regulated or business-critical environments. The center of gravity - \$20-60/hour for 1-3 months - suggests most companies treat AI agents as manageable experiments, with the potential to expand into long-term infrastructure once results are proven.



WANNA HAVE A PRECISE ESTIMATE FOR YOUR PROJECT?

The data shows a technology still taking shape, but already spreading widely. Certain tools - Python, LangChain, Pinecone - are becoming defaults, while others compete for niche roles in memory, orchestration, or voice. Use cases cluster around clear business value: automating back-office work, scaling customer support, and accelerating sales and marketing. Industry adoption follows that

logic, with sectors most focused on efficiency and customer acquisition leading the way, while others experiment at the edges. Even cost patterns reflect this duality: most projects are treated as manageable experiments, yet mission-critical builds command top-tier rates.

Together, these signals point to the same conclusion: AI agents are moving from novelty to necessity. The exact tools and applications will keep shifting, but the trajectory is clear - agents are becoming a standard layer of business operations.

This research gives a chance to see trends in developer tools for AI agents over time.

Our analysis of the AI agent development market is based on a dataset of 542 job postings collected from the freelancing platform Upwork. The data was gathered in August 2025 over a one-week period, providing a representative snapshot of the market. To provide context on our sample, we've included a breakdown of client countries below.

We analyzed several key parameters from this dataset: technologies, use cases, industries, hourly rates, and timelines. Here is a description of how each was analyzed:

- **Technologies:** We created an extensive list of relevant technologies based on a review of 100 jobs from our sample, categorized into key areas like Programming Languages, Databases and Vector Stores, LLMs and Model Providers, AI Engineering Technologies, No-code/Low-code Tools, and Voice, Speech & Audio Tech. A custom Python script was used to automatically analyze all 542 job descriptions and count the number of mentions of each technology from our list. To determine the share of each technology within its category, we used the total number of mentions for that category as the 100% baseline.
- **Use Cases & Industries:** A list of common use cases and industries was compiled from a random sample of 50 job descriptions. We then utilized a large language model (LLM) to categorize all 542 jobs based on these lists. This initial categorization was followed by a rigorous manual review and correction process to ensure accuracy and relevance.
- **Hourly Rates & Timelines:** To analyze compensation, we converted all hourly rate ranges (e.g., "\$20 - \$40") into a numerical format by averaging the minimum and maximum values. We then calculated the weighted average hourly rate for the entire sample, providing a more accurate representation of the market's standard. A similar process was applied to project timelines, where the duration (e.g., "1-3 months") was converted to a single numerical value to enable calculation of a weighted average timeline.


Read Also:

[AI Agent Development Services - Greenice](#)

The analysis of client locations reveals that the demand for AI agent development is highly concentrated in a few key countries, primarily driven by major technology and innovation hubs.

Nearly 40% of all jobs in our dataset originated from the United States, underscoring its role as the dominant client market for AI agent development. This aligns with the country's leadership in technology, venture capital, and digital transformation — and highlights how the latest technical trends in American AI agents 2026 are setting the pace for the global market. After the U.S., the largest shares come from Australia (8.7%), the United Kingdom (6.8%), and Canada (4.4%), all with strong tech ecosystems and rising demand for skilled professionals to power their digital economies.

GEOGRAPHY OF THE RESEARCH



greenice

AI agent demand is concentrated in tech hubs, with the U.S. leading at ~40%. Original research by Greenice.



READY TO BUILD YOUR AI AGENT?

Share



Inna Lebedeva

Inna Lebedeva is a market researcher and writer at Greenice web development company. She investigates IT niches and writes articles for entrepreneurs who want to launch their business in those niches. Utilizing our experienced Greenice team, and intensive market research, Inna provides in-depth analysis to business owners, enabling them to make informed decisions.



Sergii Opanasenko
Sergii is responsible for establishing and overseeing the main business operations at Greenice. In particular, he supervises web development, QA, project management, HR, and sales. Sergii's job is to ensure that all puzzle pieces of the business come together to provide the best service to our clients.



Kateryna Reshetilo
Kateryna's job is to understand the markets the company is serving, formulate the best offerings for potential clients, and market them in the most effective way. In addition to that, she is also responsible for discovering new market niches and developing competitive strategies to exploit them.

You should be
logged in
to be able to rate articles

5
★★★★★

[Login](#) to live the comment



Business Automation
Technologies

HOW TO MAKE AN AI
BOT: OWNER'S GUIDE
TO CONVERSATIONAL &
GENERATIVE AI

Envision a world where your customer service never sleeps, effortlessly handling inquiries with



Start a Project Technologies

MVP DEVELOPMENT FOR
STARTUPS

Statistics say that 90% of startups fail. Often because they try to build the "perfect" product right from the start. They try to care about the user, when they aren't sure the user exist... and has the needs the startup promises to resolve. Overcomplicating features early on



Business Automation
Technologies

BEST CRM
DEVELOPMENT
COMPANIES: TOP 10
FOR CUSTOM CRMS

Choosing the right CRM development company is one of the most important decisions for

the charm of human interaction. This isn't a distant dream—it's the reality... with AI chatbots. For business owners seeking digital elevation, knowledge acquired through the development process and a guide—more than an upgrade, it's a game-changer.

view of the pros and cons, as well as an estimate of the cost of... creating an AI chatbot.



Inna Lebedeva
06.02.2026

🕒 10 min read

drains budgets and delays market entry. We at Greenice are **masters in launching MVP (Minimal Viable Product) development** that allows you to launch a "lean, viable" version of your product, quickly, preserve resources, and ensure a solution that solves one core problem... exceptionally well. At Greenice, we believe it's not about doing less; it's about doing what matters most. It must be reliable, secure, and provide a seamless user experience. Think of it as the foundation of your future ecosystem: minimal in scope, but maximum in quality.

06.02.2026

🕒 6 min read

growing businesses. With so many options available, from boutique... agencies to enterprise businesses, we compare them to a CRM development companies know exactly how to integrate existing tools, and support long-term growth.

The table below summarizes how each company performs in key areas: security, innovation, reputation, and cost, as well as who they're best suited for. It's a quick... way to compare options, whether



Inna Lebedeva
04.02.2026

🕒 12 min read

Services

Custom Web App

Back-End Development

Custom Website

Expertise

Industries

Technologies

Company

Insights



United States
Austin, TX 78702
1201 E Cesar Chavez St



Poland
00-375 Warsaw
Juliana Bartoszewicza, 3/11



hello@greenice.net

Follow Us

Verified Reviews

©2026. Greenice

[Privacy Policy](#)

